

Formal Analysis of TPM Key Certification Protocols

©2023

Sarah Lavinia Johnson

B.S. in Computer Science, University of Kansas, 2021

B.S. in Mathematics, University of Kansas, 2021

Submitted to the graduate degree program in Department of Computer Science and the Graduate Faculty of the University of Kansas in partial fulfillment of the requirements for the degree of Master of Science in Computer Science.

Perry Alexander, Chairperson

Committee members

Michael Branicky

Emily Witt

Date defended: 03 May 2023

The Thesis Committee for Sarah Lavinia Johnson certifies
that this is the approved version of the following thesis:

Formal Analysis of TPM Key Certification Protocols

Perry Alexander, Chairperson

Date approved: TBD

Abstract

Development and deployment of trusted systems often require definitive identification of devices. A remote entity should have confidence that a device is as it claims to be. An ideal method for fulfilling this need is through the use of secure device identifiers. A secure device identifier (DevID) is defined as an identifier that is cryptographically bound to a device. A DevID must not be transferable from one device to another as that would allow distinct devices to be identified as the same. Since the Trusted Platform Module (TPM) is a secure Root of Trust for Storage, it provides the necessary protections for storing these identifiers. Consequently, the Trusted Computing Group (TCG) recommends the use of TPM keys for DevIDs. The TCG's specification *TPM 2.0 Keys for Device Identity and Attestation* describes several methods for remotely proving a key to be resident in a specific device's TPM. These methods are carefully constructed protocols which are intended to be performed by a trusted Certificate Authority (CA) in communication with a certificate-requesting device. DevID certificates produced by an OEM's CA at device manufacturing time may be used to provide definitive evidence to a remote entity that a key belongs to a specific device. Whereas DevID certificates produced by an Owner/Administrator's CA require a chain of certificates in order to verify a chain of trust to an OEM-provided root certificate. This distinction is due to the differences in the respective protocols prescribed by the TCG's specification. We aim to abstractly model these protocols and formally verify that their resulting assurances on TPM-residency do in fact hold. We choose this goal since the TCG themselves do not provide any proofs or clear justifications for how the protocols might provide these assurances. The resulting TPM-command library and execution relation modeled in Coq may easily be expanded upon to become useful in verifying a wide range of properties regarding DevIDs and TPMs.

Acknowledgements

Contents

1	Introduction	1
2	Background	2
2.1	TPM 2.0	2
2.1.1	Keys	3
2.2	Inductive Propositions	6
3	Secure Device Identity	7
3.1	Certificate Chain	9
4	Execution Model	12
5	Identity Provisioning	18
5.1	Owner/Administrator Creation of LAK Certificate based on IAK Certificate	18
5.2	OEM Creation of IAK Certificate based on EK Certificate	22
6	Conclusion	25
6.1	Conclusion	25
6.2	Future Work	25
A	Model of Inference	27
B	Model of Execution	29

List of Figures

2.1	Model of Keys	4
2.2	Model of Certificates	5
3.1	Model of Key Attribute Requirements	9
3.2	Key and Certificate Relationships [Trusted Computing Group (2021b)]	10
4.1	Model of Messages	12
4.2	Type Signature of Execute Relation	13
4.3	Model of Commands	13
4.4	Model of Command Sequences	16
4.5	Theorems on Sequential Execution	17
5.1	Parameters of LAK Certification Protocol	20
5.2	Model of Correct Steps of Owner	21
5.3	Model of Correct Steps of Owner's CA	22
5.4	Minimal Initial State of Owner	23

List of Tables

3.1	Key Requirements and Recommendations	8
-----	--	---

Chapter 1

Introduction

the following capitalized titles of TPM Manufacturer, OEM, and Owner/Administrator are keywords that will be referenced throughout this paper

simplified version of the process for creating and distributing TPM-containing devices to the end user

1. TPM Manufacturers produce TPM chips according to the international standard. They provision each TPM chip with one or more certificates which bind a key to that specific TPM. These chips are then distributed to the original equipment manufacturers (OEMs).
2. OEMs produce devices (e.g., PCs) with these TPM chips integrated. They provision each TPM chip with one or more certificates which bind a key to that specific device. These devices are then distributed to the end users (Owners/Administrators).
3. Owners/Administrators may optionally provision their TPM chip(s) with one or more certificates which bind a key to that specific device.

the provisioning of device identification in Steps 2 and 3 is the subject of this paper

Chapter 2

Background

2.1 TPM 2.0

A Trusted Platform Module (TPM) is a microcontroller that complies with the ISO/IEC 11889:2015 international standard. The TPM and its specification were designed by the Trusted Computing Group (TCG) to act as a hardware anchor for PC system security [Arthur et al. (2015)]. To this end, TPMs have the abilities necessary for secure generation of keys, algorithm agility, secure storage of keys, enhanced authorization, device health attestation, device identification, NVRAM storage and more.

The TPM's key generator is based on its own random number generator (RNG) so that it does not rely on external sources of randomness. These keys can be used for a multitude of purposes and may be created or destroyed as often as needed. Due to algorithm agility, the TPM can use nearly any cryptographic algorithm. As a result, keys may utilize asymmetric algorithms such as RSA or ECC, or they may utilize symmetric algorithms such as AES or DES. Additionally, a variety of key strengths (i.e., key sizes) and hash algorithms may be used. By design, keys stored within the TPM are protected against software attacks. Keys may optionally be further protected using enhanced authorization. Enhanced authorization (EA) allows a key or other TPM entity to be authorized using a password, HMAC, or policy. This flexibility allows for varying complexities in the requirements for accessing an entity.

Device health attestation data provided by a TPM offers cryptographic proof of software state. Attestation data comes in the form of a quote which is a signed hash over a selection of platform configuration registers. Platform configuration registers (PCRs) store the results from a chain of

boot time measurements in a way that guarantees integrity. In particular, a PCR cannot be rolled back to a previous value resulting in a measurement being undone. To tie attestation data to a specific device, the key that performed the signing operation must be cryptographically bound to that device, that is, the key must be a *secure device identifier*. There exist a variety of additional applications, such as network authentication, which similarly require a device to be definitively identified. Public key certificates may be distributed to remote entities in order to prove that the corresponding private key resides in a specific TPM-containing device. These certificates are stored within the TPM's NVRAM providing protection from accidental erasure in the scenario that the device's hard drive gets wiped.

2.1.1 Keys

Since the focus of this paper is on the usage of TPM keys for secure device identifiers, we will limit our discussion on cryptographic keys in this section to TPM keys which utilize asymmetric cryptography. A key may then be created using one of two commands.

- **TPM2_CreatePrimary:** A Primary key is produced based on the current Primary Seed. A Primary key may be persisted within the TPM. Otherwise it must be recreated after a TPM reset.
- **TPM2_Create:** An Ordinary key is produced based on a seed taken from the RNG. An Ordinary key is the child of another key; it is wrapped by that parent key. It may be persisted within the TPM or persisted external to the TPM in the form of an encrypted key blob. The blob is only loadable using the parent key's authorization in the TPM that created it.

Keys have attributes that are set at creation-time. These attributes are permanent and include the following: *FixedTPM*, *Sign*, *Decrypt*, and *Restricted*. The *FixedTPM* attribute indicates that the private key cannot be duplicated. A key pair with the *Sign* attribute set consists of a private signing key and a public signature-verification key. When properly handled, private signing keys can provide integrity, authenticity, and nonrepudiation. A key pair with the *Decrypt* attribute set

consists of a public encryption key and a private decryption key. When properly handled, public encryption keys can provide confidentiality. A key with both the Sign and Decrypt attributes set is called a Combined key. US NIST SP800-57 disallows the use of Combined keys for the reason that it may weaken the security guarantees associated with one or both of the attributes. Furthermore, a key pair may have the Restricted attribute set, limiting the operations of the private key to TPM generated data.

The Coq model inductively defines a `pubKey` and `privKey` type for public keys and private keys respectively. A key of either of these types requires a unique identifier and a sequence of boolean values describing whether a particular attribute is set or not set. A key pair consists of a `pubKey` and a `privKey` with the same identifier and attributes. The model does not differentiate between Primary and Ordinary keys.

```
Inductive pubKey : Type :=
| Public : keyIdType → Restricted → Sign → Decrypt → FixedTPM → pubKey.

Inductive privKey : Type :=
| Private : keyIdType → Restricted → Sign → Decrypt → FixedTPM → privKey.
```

Figure 2.1: Model of Keys

The Restricted attribute provides important security implications. A restricted signing key may only sign a digest that has been produced by the TPM. Enforcement of this constraint is reliant on a 4-byte magic value called `TPM_Generated` [Trusted Computing Group (2019)]. All structures that the TPM constructs from internal data begins with this value. Such structures include keys and PCRs. These structures result in several interesting and significant uses for restricted signing keys, namely key certification and attestation. In particular, a restricted signing key is used during key certification to prove that a new key resides in the same TPM as itself. And during attestation, a restricted signing key is used to prove that a quote is the result of the PCR values within the same TPM in which the key itself resides. Additionally, a restricted signing key has the ability to sign data supplied to the TPM externally by using the `TPM2_Hash` command. In this case, the `TPM2_Hash` command produces a ticket asserting that the TPM itself calculated this hash and will

later sign it. A restricted signing key will not sign external data without this ticket. To prevent spoofing of another TPM’s internal data as one’s own, the TPM2_Hash command will only produce a ticket if the external data does not begin with the TPM_Generated value.

A restricted decryption key is called a storage key. Only storage keys can be used as parents to create or load child objects or to activate credentials [Arthur et al. (2015)]. All TPMs are shipped with an essential storage key: the endorsement key. The endorsement key (EK) is installed by the TPM Manufacturer and stored in a shielded location on the TPM. The corresponding EK certificate serves a significant role in the enrollment of secure device identifiers. This process will be discussed in further detail in later sections.

A certificate contains a public key and an identity and is signed by a trusted Certificate Authority. A certificate binds a public key to an identity. The term certificate specifically refers to an X.509 v3 digital certificate. The EK certificate includes the public part of the EK itself as well as various assertions regarding the security qualities and provenance of the TPM [Trusted Computing Group (2021a)]. The EK certificate binds the EK to a specific TPM. For keys created by entities other than the TPM manufacturer (i.e., the OEM and the Owner/Administrator), a certificate’s identity field will contain non-TPM device information. This information should be globally unique per device [Institute of Electrical and Electronics Engineers (2018)]. In this model, certificates are abstractly defined as the inductive signedCert type. A signedCert requires a public key, an identifier, and a private key. An identifier may include information describing either the TPM or the device. The private key parameter denotes the key which performed the signature over the certificate.

```
Inductive identifier : Type :=
| TPM_info : tpmInfoType → identifier
| Device_info : deviceInfoType → identifier.

Inductive signedCert : Type :=
| Cert : pubKey → identifier → privKey → signedCert.
```

Figure 2.2: Model of Certificates

2.2 Inductive Propositions

[Pierce et al. (2022)]

Chapter 3

Secure Device Identity

A secure device identifier (DevID) is defined as an identifier that is cryptographically bound to a device [Institute of Electrical and Electronics Engineers (2018)]. A device with DevID capability includes at least one Initial Device Identifier (IDevID). An IDevID must not be transferable from one device to another and must be stored in a way that protects it from modification. Furthermore, an IDevID is intended to be long-lived and usable for the lifetime of the device. Since the TPM is a secure Root of Trust for Storage and protects keys against compromise, TPM keys are an ideal choice for IDevIDs. The TCG specifically makes this recommendation as well. IDevIDs are installed by OEMs in TPM-containing devices at manufacturing time. Additionally, a device with TPM-based DevID capability may support the creation of Locally Significant Device Identifiers (LDevIDs) by the end user (i.e., the Owner/Administrator). An LDevID must not be transferable to a device with a different IDevID without knowledge of the private key used to produce the cryptographic binding.

When using a TPM key for secure device identity, there are restrictions on the attributes that it can have in order to enforce the best security practices; the key must have the FixedTPM and Sign attributes set and the Decrypt attribute not set. The Restricted attribute may optionally be set. When the Restricted attribute is set, such a key is called an attestation key (AK). This DevID gets its special name due to its unique ability to be used as a parent node in a chain of certificates. We will discuss this idea in later sections. The acronym AK is prefixed by the letter I or L denoting Initial or Locally Significant respectively. The FixedTPM attribute must always be set because it is of paramount importance that a DevID never be duplicated, transferred, or copied. OEM-installed DevIDs (i.e., IAKs and IDevIDs) should be Primary keys so that they may be recoverable by the

Key	Type	FixedTPM	Signing	Decrypting	Restricted	Creator
EK	Primary	X		X	X	TPM Manufacturer
IAK	Primary	X	X		X	OEM
IDevID	Primary	X	X			OEM
LAK	Ordinary	X	X		X	Owner/Admin
LDevID	Ordinary	X	X			Owner/Admin

Table 3.1: Key Requirements and Recommendations

Owner/Administrator. Since an Owner/Administrator cannot provision new IAKs or IDevIDs on their device, these keys should be able to be recreated during the lifetime of the TPM or more precisely during the lifetime of the TPM’s Primary Seed to avoid the problematic loss of these essential identifiers. An IAK is intended to be used only for the certification of new DevIDs. On the other hand, an LAK may be used for attestation as well as for the certification of new DevIDs. An IDevID or LDevID may be used for general device authentication purposes.

The issuers of device identity certificates are known as Certificate Authorities (CAs). CAs are further identified by the creator of the keys that they certify (i.e., the CA that issues certificates for IAKs and IDevIDs is known as the OEM’s CA and the CA that issues certificates for LAKs and LDevIDs is known as the Owner/Administrator’s CA). The OEM’s CA must carefully verify the attributes and TPM residency of a key before signing a certificate due to the important security and identity implications provided by these certificates. The Owner/Administrator’s CA should ideally also verify the attributes and TPM residency of a key before signing a certificate. The attribute requirements displayed in Table 3.1 is modeled as a collection of functions. Each function takes a public key as input and returns a proposition. If a CA issues a certificate, then it should be the case that applying the corresponding function on the DevID results in the value True. All CAs should support a standard certificate transport protocol that provides confidentiality, integrity, and protection from replay attacks [Trusted Computing Group (2021b)]. These transport protocols are outside the scope of this paper. We will assume CAs to be following this recommendation precisely.

```

Definition endorsementKey (k : pubKey) : Prop :=
  match k with
  | Public _ Restricting NonSigning Decrypting Fixing  $\Rightarrow$  True
  | _  $\Rightarrow$  False
  end.

Definition attestationKey (k : pubKey) : Prop :=
  match k with
  | Public _ Restricting Signing NonDecrypting Fixing  $\Rightarrow$  True
  | _  $\Rightarrow$  False
  end.

Definition devidKey (k : pubKey) : Prop :=
  match k with
  | Public _ NonRestricting Signing NonDecrypting Fixing  $\Rightarrow$  True
  | _  $\Rightarrow$  False
  end.

```

Figure 3.1: Model of Key Attribute Requirements

3.1 Certificate Chain

A chain of certificates can be used to verify a chain of trust to some trust anchor. The IAK certificate typically acts as this trust anchor. In issuing an IAK certificate, the OEM’s CA makes an assertion that is a primary security dependency for future enrollment of all DevIDs [Trusted Computing Group (2021b)]. Since only AKs (i.e., DevIDs with the Restricted attribute set) may be used for key certification, enrollment of all DevIDs is reliant, either directly or indirectly, on an IAK certificate. We can see in Figure 3.2 that the IDevID, LAK, and LDevID may all be linked back to the IAK. The creation of an IAK certificate relies on the EK certificate. The TPM Manufacturer’s method for creation of an EK certificate is outside the scope of this paper since an EK is not a DevID. We trust that an EK certificate provides definitive evidence that the EK resides within the specific TPM.

IAK certificates provide definitive evidence to a remote entity that a key belongs to a specific device. To demonstrate that some other key belongs to a specific device, one must provide all of the certificates which link that key to an IAK. These dependencies are a direct result of the protocols performed to issue these certificates.

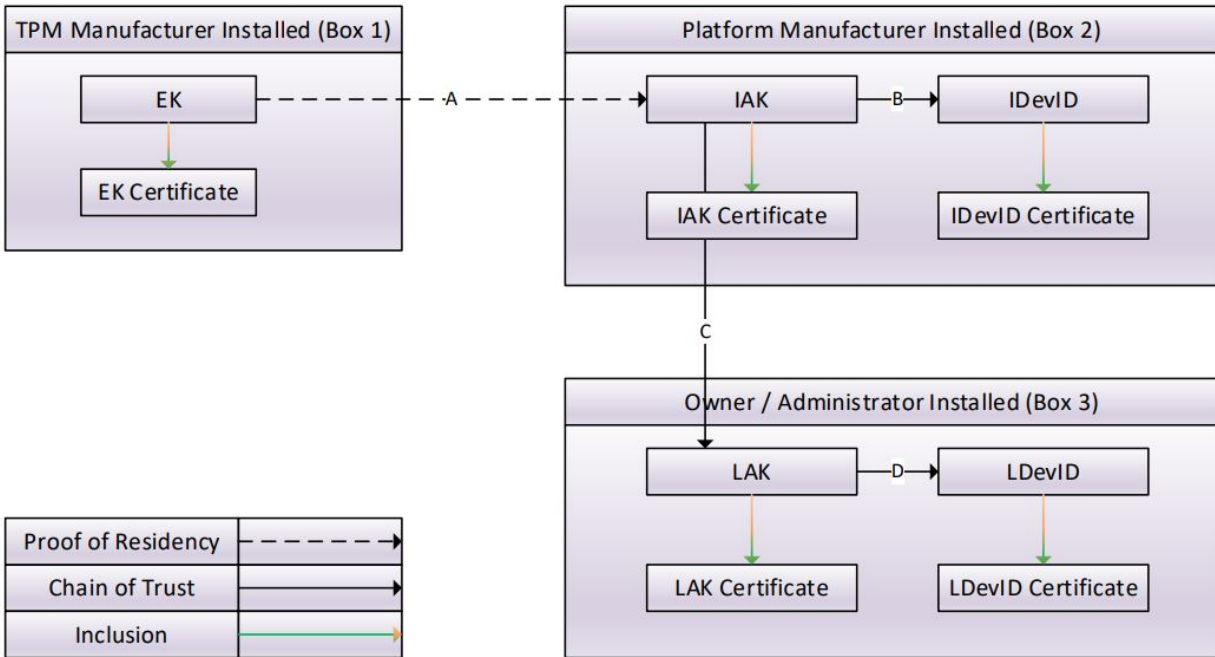


Figure 3.2: Key and Certificate Relationships [Trusted Computing Group (2021b)]

- Box 1: The EK certificate is signed by the TPM Manufacturer's CA and binds the EK to a specific TPM.
- Line A: The IAK is verified by the OEM's CA to have the correct key properties and to be resident in the same TPM as the EK.
- Line B: The IDevID is verified by the OEM's CA to have the correct key properties and to be resident in the same TPM as the IAK.
- Box 2: The IAK certificate and IDevID certificate is signed by the OEM's CA and binds the IAK and IDevID to a specific device.
- Line C: The LAK is verified by the Owner/Administrator's CA to have the correct key properties and to be resident in the same TPM as the IAK.
- Line D: The LDevID is verified by the Owner/Administrator's CA to have the correct key properties and to be resident in the same TPM as the LAK.
- Box 3: The LAK certificate and LDevID certificate is signed by the Owner/Administrator's CA.

Figure 3.2 shows the relationships between keys and certificates for exactly one of each type of DevID. In practice, there often exists multiple of each type of DevID. Due to algorithm agility, an OEM usually installs multiple IAKs and IDevIDs with varying algorithms and sizes. An Owner/Administrator may install new LAKs and LDevIDs as often as need. While an IDevID certificate always relies directly on an IAK certificate, an LAK or LDevID certificate may be directly reliant on either an IAK or LAK certificate. In particular, an LAK certificate may be used to enroll LDevIDs or additional LAKs creating an arbitrary long chain of certificates. All DevIDs (excluding IAKs) must be directly linked to an AK because the Restricted attribute is necessary in proving that the new key resides in a specific TPM (i.e., the same TPM as the AK).

Chapter 4

Execution Model

The protocols used to enroll DevID certificates require both TPM and non-TPM commands to be performed by the CA and the requesting device. A command may rely on a variety of parameters such as keys, nonces, certificates, as well as other messages. A message includes all of the structures that an entity may use or produce. The message type is an abstract representation of these

```
Inductive message : Type :=  
| publicKey : pubKey → message  
| privateKey : privKey → message  
| hash : message → message  
| signature : message → privKey → message  
| TPM2B_Attest : pubKey → message  
| encryptedCredential : message → randType → pubKey → message  
| randomNum : randType → message  
| TCG_CSR_IDevID : identifier → signedCert → pubKey → message  
| TCG_CSR_LDevID : message → signedCert → message  
| signedCertificate : signedCert → message  
| pair : message → message → message.
```

Figure 4.1: Model of Messages

structures. From a message, additional messages may be inferred. For example, given a message in the form `signature m k`, the message `m` may be deduced. Whereas given a message in the form `encryptedCredential n g k`, no messages may be deduced. This concept is modeled in two ways: as a recursive function `inferFrom` and as an inductive proposition `inferrable`. These two definitions are proven to be equivalent (see Appendix A for definitions and proof). In particular, additional information may be gained from signatures, TPM2B_Attest structures, certificate signing requests (CSRs), public key certificates, and pairs of messages. All other messages either

contain no additional information (i.e., keys and random numbers) or the information is concealed (i.e., hash digests and encryptions).

Each command and its execution is modeled abstractly. We do not attempt to model the computational intricacies of true cryptography. Command execution is defined as an inductive proposition relating an initial state pair, a command, and a final state pair (see Appendix B for definition). The

Inductive `execute : tpm_state * state → command → tpm_state * state → Prop`

Figure 4.2: Type Signature of Execute Relation

`tpm_state` and `state` types are aliases for a list of messages. These types are implemented as a list only for convenience; they are treated as a set in all practical aspects (i.e., ordering and duplicates are ignored). The distinction between these two state types will become clear as we inspect each command and its corresponding real-life and modeled execution behaviors. Each command

Inductive `command : Type :=`
`| CheckAttributes : pubKey → Restricted → Sign → Decrypt → command`
`| TPM2_Hash : message → command`
`| CheckHash : message → message → command`
`| TPM2_Sign : message → privKey → command`
`| TPM2_Certify : pubKey → privKey → command`
`| CheckSig : message → pubKey → command`
`| MakeCSR_IDevID : identifier → signedCert → pubKey → command`
`| MakeCSR_LDevID : message → signedCert → command`
`| CheckCert : signedCert → pubKey → command`
`| MakePair : message → message → command`
`| CheckRandom : message → randType → command`
`| TPM2_ActivateCredential : message → privKey → privKey → command.`

Figure 4.3: Model of Commands

corresponds with exactly one constructor in the `execute` relation (excluding `TPM2_Sign` which corresponds with two). Each constructor possesses its own distinctive conditions; these conditions must be met for execution to be considered successful. These conditions are manifested in several forms. For one, constructors pattern match on a command's inputs. Additionally, certain messages may be required to be in the initial `tpm_state` or `state`. And lastly, the command's results must

be added to the final `tpm_state` or `state`.

The `CheckAttributes` command verifies that a public key has all of the provided attributes. In practice, this is done by checking the `TPMA_Object` bits of the key. In the model, these values are stored within the `Restricted`, `Sign`, `Decrypt`, and `FixedTPM` fields of the `pubKey` type. In order to check the attributes of a particular key, one must have knowledge of that key. The key need not reside within one's own TPM though since this command is typically used to check the attributes of some external entity's key. The model requires specifically that the initial state contains the public key to be checked.

The `TPM2_Hash` command performs a cryptographic hash operation on a piece of data. This data may be any message that is known to the entity performing the command; this message must be contained in the initial state. The result of the operation is abstractly defined using the opaque hash constructor and is stored in both the final `tpm_state` and `state`. Only one command may be used to determine the contents of a hash digest, that is, the `CheckHash` command which verifies that the contents of a hash digest match a particular plaintext message. Both the hash digest and the plaintext message must be contained in the initial state.

PCRs are not currently included in the message language

The `TPM2_Sign` command generates a signature over a message using the specified private key. There are several conditions for successful execution. For one, the key must have the `Sign` attribute set. Additionally, the key must reside in the TPM (i.e., be contained in the initial `tpm_state`).
different based on restricted The execute relations If the key has the `Restricted` attribute not set, then the message must simply be in the initial state. On the other hand, if the key does have the `Restricted` attribute set, then the message must have been produced by the TPM and be in the TPM state. As described in Section 2.1.1, a restricted signing key may sign external data if and only if the TPM first performed a hash operation on the data. We call the hash operation a signature hash when used for this purpose. In practice, the `TPM2_Hash` command would produce a ticket containing a validation structure which indicates that the resulting hash was produced by the TPM and is safe to sign. This ticket is then passed to the `TPM2_Sign` command. In the model,

these tickets are handled implicitly: the hash digest produced by TPM2_Hash is added to the TPM state so that the restricting sign key may then sign the external data.

The TPM2_Certify command proves that an object is loaded in the TPM by producing a signed TPM2B_Attest structure. The command requires two inputs: a public key to be certified and a private key to sign the attestation structure. The private key must have the Sign attribute set and must reside in the TPM. Upon receiving a request to execute the TPM2_Certify command, the TPM will verify that the inverse of the public key parameter resides in the TPM as well. Messages produced by the TPM2_Sign and TPM2_Certify commands are defined using the signature constructor. A signature may be verified against a public key using the CheckSig command. If the provided public key is the inverse of the private key which performed the signature, then the check succeeds.

The MakeCSR_IDevID command produces a TCG_CSR_IDevID structure containing the provided inputs. A TCG_CSR_IDevID is a certificate signing request (CSR) which contains the data required to couple an IAK to a TPM-containing device. Additionally, it may include the certification information for an IDevID if one wishes to produce both the IAK and IDevID certificates in a single pass. In particular, this structure is used any time an enrollment process uses the EK certificate. The Trusted Computing Group (TCG) defines a C typedef structure to group all of the fields. In the model, we only include the fields necessary for creating an IAK certificate (i.e., device-identifying information, the EK certificate, and a public key to be certified). The MakeCSR_LDevID command is very similar to the MakeCSR_IDevID command except it produces a TCG_CSR_LDevID structure which includes the certification information for an LAK or LDevID. In the model, we only include the fields necessary for creating an LAK certificate (i.e., a signed TPM2B_Attest structure and the IAK certificate).

The CheckCert command verifies a signature over a certificate against a public key. One should check an EK certificate against the public key of the TPM Manufacturer's CA, an IAK or IDevID certificate against the public key of the OEM's CA, and an LAK or LDevID certificate against the public key of the Owner/Administrator's CA.

The `MakePair` command combines two messages into a single message using the `pair` constructor.

Due to the abstract, symbolic nature of this model, several TPM commands are intentionally excluded from the `command` type. This results specifically from an inability to truly capture the cryptographic properties of randomness. Randomness plays a vital role in the real-life implementation of keys and nonces, that is, it prevents a key or nonce from being guessed. Since we are unable to preserve this property in our model, we choose to eliminate all commands which generate a key or nonce. Therefore, a message of either of these types must be inferred from some other message or be in the initial state.

The `TPM2_GetRandom` command retrieves random bytes from the TPM. These resulting bytes may be used as a nonce. Due to the reason described above, we do not include this command explicitly in our model, although we do assume it is used by the CA in some protocols.

The `TPM2_MakeCredential` command

The `TPM2_ActivateCredential` command

Commands in the model are sequenced linearly by the `sequence` type. In fact, this type is identical in structure to the Coq type `list command`. Sequential command execution is defined as an inductive proposition relating an initial state, a command sequence, and a final state. We

```

Inductive sequence : Type :=
| Sequence : command → sequence → sequence
| Done : sequence.

Infix ";;" := Sequence (at level 60, right associativity).

Inductive seq_execute : tpm_state * state → sequence → tpm_state * state → Prop :=
| SE_Seq : ∀ ini mid fin c s,
  execute ini c mid →
  seq_execute mid s fin →
  seq_execute ini (Sequence c s) fin
| SE_Done : ∀ ini,
  seq_execute ini Done ini.

```

Figure 4.4: Model of Command Sequences

can prove several interesting and useful facts about the `seq_execute` relation. Firstly, sequential execution is deterministic. Given an initial state and a command sequence, there is at most one final state which satisfies the `seq_execute` relation. This means that `seq_execute` is a partial function. Next, sequential execution is an expansion. Given a related initial state, command sequence, and final state, the initial state is always a subset of the final state (i.e., `seq_execute` expands on the initial state). In particular, commands do not remove elements from the state. Finally, sequential execution cannot generate keys or nonces. This feature is due to the deliberate exclusion of certain TPM commands from the language. The proof statements are recorded below (See Appendix B for the proofs themselves).

Theorem `seq_exec_deterministic` : $\forall \text{ini } s \text{ fin1 fin2},$
 $\text{seq_execute } \text{ini } s \text{ fin1} \rightarrow$
 $\text{seq_execute } \text{ini } s \text{ fin2} \rightarrow$
 $\text{fin1} = \text{fin2}.$

Theorem `seq_exec_expansion` : $\forall \text{iniTPM } \text{ini } s \text{ finTPM } \text{fin},$
 $\text{seq_execute } (\text{iniTPM}, \text{ini}) \text{ } s \text{ } (\text{finTPM}, \text{fin}) \rightarrow$
 $(\text{iniTPM} \subseteq \text{finTPM}) \wedge (\text{ini} \subseteq \text{fin}).$

Figure 4.5: Theorems on Sequential Execution

Chapter 5

Identity Provisioning

To maintain a cryptographic evidentiary chain linking a DevID to a specific TPM and device, the CA should follow certain provisioning protocols. The TCG describes several such protocols in their specification *TPM 2.0 Keys for Device Identity and Attestation*. We will consider in detail two of these protocols: OEM creation of an IAK certificate based on an EK certificate and Owner/Administrator creation of an LAK certificate based on an IAK certificate. We choose these two protocols since they bear the most significance in enrollment of additional DevIDs (recall that AK certificates may be used as parent nodes in a chain of certificates). For each protocol, the TCG's specification not only outlines its steps but also claims it provides certain assurances.

For convenience and clarity, we will inspect each of these protocols in the reverse order that their dependencies entails.

5.1 Owner/Administrator Creation of LAK Certificate based on IAK Certificate

In this section, we will shorten the term Owner/Administrator to Owner. The Administrator may still be a participant in this protocol. We only mean for this to act as an abbreviation. Now, we begin by provided a description of the recommended procedure.

0. The Owner creates and loads the LAK
1. The Owner certifies the LAK with the IAK
2. The Owner builds the CSR containing:
 - (a) The signed TPM2B_Attest structure

- (b) The IAK certificate
- 3. The Owner takes a signature hash of the CSR
- 4. The Owner signs the resulting hash digest with the LAK
- 5. The Owner sends the CSR paired with the signed hash to the CA
- 6. The CA verifies the recieved data by checking:
 - (a) The hash digest against the CSR
 - (b) The signature on the hash digest with the LAK public key
 - (c) The signature on the TPM2B_Attest structure with the IAK public key
 - (d) The signature on the IAK certificate with the public key of the OEM's CA
 - (e) The attributes of the LAK
- 7. If all of the checks succeed, the CA issues the LAK certificate to the Owner

The TCG's specification claims that this procedure provides the following assurances: (A) the new LAK is resident in the same TPM as the IAK and (B) the LAK has the correct key properties. No proof or clear justification is presented to support this claim though. Our goal is to prove that these assurances do in fact hold.

We model this protocol within Coq's `Module Type` mechanism. This mechanism allows for the inclusion of parameters which provides the necessary flexibility to describe the protocol generally. Each participating entity has its own parameters. The Owner has a key to be certified (i.e., the LAK), its IAK, and IAK certificate. The CA has its own key and the public key of the OEM's CA. The `Module Type` mechanism additionally allows for axioms to be defined. When instantiating a `Module Type` with concrete values, one must prove all of the axioms. The first axiom we define is straightforward and only attempts to enforce the randomness of cryptographic keys, that is, all key parameters must be pairwise distinct.

The procedure may be regarded as the composition of two sequences: the Owner's steps (i.e., steps 0-5) followed by the CA's steps (i.e., steps 6-7). With that in mind, we can construct an object of type `sequence` for the Owner. Since the `command` type does not include a method for creating and loading keys, we assume step 0 to have been performed prior to the sequence. The CA's

```

(* Owner parameters *)
Parameter pubLAK : pubKey.
Parameter pubIAK : pubKey.
Parameter certIAK : signedCert.

(* CA parameters *)
Parameter pubCA : pubKey.
Parameter pubOEM : pubKey.

(* All keys are pairwise distinct *)
Axiom keys_distinct :
  pubLAK <> pubIAK ∧
  pubLAK <> pubCA ∧
  pubLAK <> pubOEM ∧
  pubIAK <> pubCA ∧
  pubIAK <> pubOEM ∧
  pubCA <> pubOEM.

```

Figure 5.1: Parameters of LAK Certification Protocol

steps are more complicated and are constructed as a function type. First the CA waits to receive a certification request from the Owner (see the message parameter `msg`). The request must be in a specific format to be considered valid (see the match statement on `msg`). The CA then executes step 6 of the procedure (see the sequence within `seq_execute`). If execution succeeds, the CA will issue the LAK certificate to the Owner (see the `Prop` return type). We include several additional parameters and criteria to serve as a method for referencing certain elements of the certification request within proof statements.

Going forward, we will consider two cases: (1) the Owner and the CA are both trusted to execute their steps correctly and (2) only the CA is trusted to execute its steps correctly. The TCG's specification does not state which of these assumptions they reason under.

Let us first consider case 1: the Owner and the CA are both trusted to execute their steps correctly. We will begin by examining the Owner's steps and its requirements. We can quantitatively describe these requirements by a minimal initial state pair. Given a sequence, a minimal initial state pair is defined as the smallest TPM state and general state which allows for successful execution of the sequence. Note that this definition does not constrain the general state to be a superset of the

```

Definition steps1to5_Owner : sequence :=
  TPM2_Certify
    pubLAK
    privIAK ;;
  MakeCSR_LDevID
    (signature (TPM2B_Attest pubLAK) privIAK)
    certIAK ;;
  TPM2_Hash
    (TCG_CSR_LDevID (signature (TPM2B_Attest pubLAK) privIAK) certIAK) ;;
  TPM2_Sign
    (hash (TCG_CSR_LDevID (signature (TPM2B_Attest pubLAK) privIAK) certIAK))
    privLAK ;;
  MakePair
    (TCG_CSR_LDevID (signature (TPM2B_Attest pubLAK) privIAK) certIAK)
    (signature (hash (TCG_CSR_LDevID (signature (TPM2B_Attest pubLAK) privIAK) certIAK)) privLAK) ;;
  Done.

```

Figure 5.2: Model of Correct Steps of Owner

TPM state; in fact, it is the case that for most sequences it is not. We intuitively construct a minimal initial state pair for `steps1to5_Owner`. First, the private LAK and private IAK must reside in the same TPM because the Owner certifies the LAK with the IAK. Next, the IAK certificate must be in state because the Owner includes it in the CSR. This intuition is useful in proving that this resulting state pair is in fact minimal. The proof statement is constructed by two parts. First, the minimal initial state is a lower bound on the set of possible initial states. And second, the minimal initial state is sufficient for successful execution.

The lower bound property is proven to hold in general. On the other hand, the sufficiency property requires that the IAK and LAK have "good" attributes. Specifically both keys must have the Sign attribute set in order to execute the `TPM2_Certify` and `TPM2_Sign` commands respectively. We choose to make the broader assumption that the IAK and LAK have all of the attributes associated with a good attestation key. See Appendix x for the proofs of these two statements.

```

Definition steps_CA (msg : message) (iak lak : pubKey) (cert : signedCert) : Prop :=
  match msg with
  | (pair (TCG_CSR_LDevID (signature (TPM2B_Attest k) k0')
    (Cert k0 id k_ca'))
    (signature m k')) =>
    iak = k0 ∧
    lak = k ∧
    cert = (Cert k0 id k_ca') ∧
    seq_execute (iniTPM_CA, inferFrom m ++ ini_CA)
      (CheckHash m
        (TCG_CSR_LDevID (signature (TPM2B_Attest k) k0')
          (Cert k0 id k_ca')) ;;
        CheckSig (signature m k') k ;;
        CheckSig (signature (TPM2B_Attest k) k0') k0 ;;
        CheckCert (Cert k0 id k_ca') pubOEM ;;
        CheckAttributes k Restricting Signing NonDecrypting ;;
        Done)
      (iniTPM_CA,
        inferFrom m ++ ini_CA)
  | _ => False
end.

```

Figure 5.3: Model of Correct Steps of Owner's CA

5.2 OEM Creation of IAK Certificate based on EK Certificate

0. The OEM creates and loads the IAK
1. The OEM builds the CSR containing:
 - (a) Device identity information including the device model and serial number
 - (b) The EK certificate
 - (c) The IAK public area
2. The OEM takes a signature hash of the CSR
3. The OEM signs the resulting hash digest with the IAK
4. The OEM sends the CSR paired with the signed hash to the CA
5. The CA verifies the recieved data by checking:
 - (a) The hash digest against the CSR
 - (b) The signature on the hash digest with the IAK public key

```

Definition iniTPM_Owner : tpm_state :=
[ privateKey privLAK ;
  privateKey privIAK ].

Definition ini_Owner : state :=
[ signedCertificate certIAK ].

Lemma ini_Owner_lowerBound :  $\forall$  iniTPM ini fin,
  seq_execute (iniTPM, ini) steps1to5_Owner fin  $\rightarrow$ 
  (iniTPM_Owner  $\subseteq$  iniTPM)  $\wedge$ 
  (ini_Owner  $\subseteq$  ini).

Lemma ini_Owner_sufficient :  $\forall$  m,
  attestationKey pubIAK  $\rightarrow$ 
  steps_CA m pubIAK pubLAK certIAK  $\rightarrow$ 
   $\exists$  fin, seq_execute (iniTPM_Owner, ini_Owner) steps1to5_Owner fin.

```

Figure 5.4: Minimal Initial State of Owner

- (c) The signature on the EK certificate with the public key of the TPM Manufacturer's CA
- (d) The attributes of the IAK
- 6. If all of the checks succeed, the CA issues a challenge blob to the OEM by:
 - (a) Calculating the cryptographic name of the IAK
 - (b) Generating a nonce
 - (c) Building the encrypted credential structure using the name of the IAK, the nonce, and the EK public key
- 7. The OEM releases the secret nonce by verifying the name of the IAK and decrypting the challenge blob
- 8. The CA checks the returned nonce against the one generated in step 6b
- 9. If the check succeeds, the CA issues the IAK certificate to the OEM

We model this protocol within Coq's Module Type mechanism. This mechanism allows for the inclusion of parameters which provides the necessary flexibility to describe the protocol in the general case. Each participating entity has its own parameters. The OEM has a key to be certified (i.e., the IAK), its EK certificate, and information to identify its device. The CA has its own key, the public key of the TPM Manufacturer's CA, and a nonce. The nonce parameter may be re-

moved if we include the `TPM2_GetRandom` command in the model's language. The `Module Type` mechanism additionally allows for axioms to be defined. When instantiated a `Module Type` with concrete values, one must prove all of the axioms. The first axiom we define is straightforward and only attempts to enforce the randomness of cryptographic keys, that is, all key parameters must be pairwise distinct.

Chapter 6

Conclusion

6.1 Conclusion

6.2 Future Work

Include attestation variety of protocols which uses PCRs.

Include more TPM and TSS commands in the model to create a library. Find a way to add TPM2_Create and TPM2_GetRandom.

Include additional control sequences such as command branching, looping, etc

Publish at iFM.

References

- Arthur, W., Challener, D., & Goldman, K. (2015). *A Practical Guide to TPM 2.0: Using the Trusted Platform Module in the New Age of Security*. Apress.
- Institute of Electrical and Electronics Engineers (2018). IEEE Standard for Local and Metropolitan Area Networks - Secure Device Identity. *IEEE Std 802.1AR-2018*.
- Pierce, B. C., de Amorim, A. A., Casinghino, C., Gaboardi, M., Greenberg, M., Hrițcu, C., Sjöberg, V., & Yorgey, B. (2022). *Logical Foundations*. Software Foundations series, volume 1. Electronic textbook.
- Trusted Computing Group (2019). *Trusted Platform Module Library Specification, Family 2.0, Level 00, Revision 01.59*.
- Trusted Computing Group (2021a). *TCG EK Credential Profile For TPM Family 2.0, Level 00, Version 2.4, Revision 3*.
- Trusted Computing Group (2021b). *TPM 2.0 Keys for Device Identity and Attestation*.

Appendix A

Model of Inference

```
Lemma inferFrom_iff_inferrable :  $\forall$  m st,  
  inferFrom m = st  $\leftrightarrow$  inferrable m st.
```

Proof.

```
  intros m st; split; intros H.  
  - generalize dependent m; assert (HI :  $\forall$  m, inferrable m (inferFrom m)); intros m.  
  -- induction m; simpl; try destruct c; try destruct s;  
    repeat constructor; assumption.  
  -- intros H; induction m; subst; apply HI.  
  - induction H; simpl; subst; try destruct crt; reflexivity.
```

Qed.

```
Fixpoint inferFrom (msg : message) : state :=  
  match msg with  
  | signature m k  $\Rightarrow$   
    (signature m k :: inferFrom m)  
  | TPM2B_Attest k  $\Rightarrow$   
    [TPM2B_Attest k ; publicKey k]  
  | TCG_CSR_IDevID id0 (Cert k id k_ca) k0  $\Rightarrow$   
    [TCG_CSR_IDevID id0 (Cert k id k_ca) k0 ; publicKey k0 ;  
     signedCertificate (Cert k id k_ca) ; publicKey k]  
  | TCG_CSR_LDevID m (Cert k id k_ca)  $\Rightarrow$   
    (TCG_CSR_LDevID m (Cert k id k_ca) :: inferFrom m ++  
     [signedCertificate (Cert k id k_ca) ; publicKey k])  
  | signedCertificate (Cert k id k_ca)  $\Rightarrow$   
    [signedCertificate (Cert k id k_ca) ; publicKey k]  
  | pair m1 m2  $\Rightarrow$   
    (pair m1 m2 :: inferFrom m1 ++ inferFrom m2)  
  | _  $\Rightarrow$   
    [msg]  
  end.
```

```

Inductive inferrable : message → state → Prop :=
| I_publicKey : ∀ k,
  inferrable (publicKey k)
    [publicKey k]
| I_privateKey : ∀ k,
  inferrable (privateKey k)
    [privateKey k]
| I_hash : ∀ m,
  inferrable (hash m)
    [hash m]
| I_signature : ∀ m k st,
  inferrable m st →
  inferrable (signature m k)
    (signature m k :: st)
| I_Attest : ∀ k,
  inferrable (TPM2B_Attest k)
    [TPM2B_Attest k ; publicKey k]
| I_encryptedCredential : ∀ n g k,
  inferrable (encryptedCredential n g k)
    [encryptedCredential n g k]
| I_randomNum : ∀ g,
  inferrable (randomNum g)
    [randomNum g]
| I_CSR_IDevID : ∀ id crt k st,
  inferrable (signedCertificate crt) st →
  inferrable (TCG_CSR_IDevID id crt k)
    (TCG_CSR_IDevID id crt k :: publicKey k :: st)
| I_CSR_LDevID : ∀ m crt st1 st2,
  inferrable m st1 →
  inferrable (signedCertificate crt) st2 →
  inferrable (TCG_CSR_LDevID m crt)
    (TCG_CSR_LDevID m crt :: st1 ++ st2)
| I_signedCertificate : ∀ k id k_ca,
  inferrable (signedCertificate (Cert k id k_ca))
    [signedCertificate (Cert k id k_ca) ; publicKey k]
| I_pair : ∀ m1 m2 st1 st2,
  inferrable m1 st1 →
  inferrable m2 st2 →
  inferrable (pair m1 m2)
    (pair m1 m2 :: st1 ++ st2).

```

Appendix B

Model of Execution

```
Inductive execute : tpm_state * state → command → tpm_state * state → Prop :=
| E_CheckAttributes : ∀ stTPM st i r s d f,
  In (publicKey (Public i r s d f)) st →
  execute (stTPM, st)
    (CheckAttributes (Public i r s d f) r s d f)
    (stTPM, st)
| E_Hash : ∀ stTPM st m,
  In m st →
  execute (stTPM, st)
    (TPM2_Hash m)
    (hash m :: stTPM, hash m :: st)
| E_CheckHash : ∀ stTPM st m,
  In (hash m) st →
  In m st →
  execute (stTPM, st)
    (CheckHash (hash m) m)
    (stTPM, st)
```